

Database fundamentals summary (Part 3)

Agenda

1. SQL Subqueries (Nested Queries)

- Basic Subquery
- Using `ALL`
- Using `ANY` / `SOME`

2. SQL Aggregate Functions

- Common functions: `MAX`, `MIN`, `AVG`, `SUM`, `COUNT`
- Practical examples

3. `GROUP BY` & `HAVING`

4. SQL Views

- Definition & Use Cases
- Creating a View
- Querying Views
- Updating / Modifying Views
- Deleting Views
- Advanced Options (`WITH CHECK OPTION`, Permissions)
- Advantages of Views

5. SQL Index

- What is an Index?
- Advantages & Disadvantages
- Best Practices
- Syntax & Examples

- Key Idea: **Speed vs Overhead**

SQL Subqueries (Nested Queries)

1. Basic Subquery

- A query **inside another query**.
- **Used to compare against a single computed value.**

```
SELECT fname, salary
-- Outer query finds all employees earning more than that salary.
FROM Employee
WHERE salary > (
    SELECT salary -- Inner query finds Ahmed Ali's salary.
    FROM Employee WHERE fname = 'Ahmed' AND lname = 'Ali'
);
```

2. Using **ALL**

- Condition must be true **against ALL values** returned by subquery.
- If any value fails, row is excluded.

```
SELECT fname, salary FROM Employee
-- Employee salary must be higher than every salary in dept 10.
WHERE salary > ALL (
    SELECT salary FROM Employee WHERE dno = 10
);
```

3. Using **ANY** / **SOME**

- Condition must be true for **at least one value** from subquery.
- More flexible than **ALL**.

```
SELECT fname, salary FROM Employee
-- Employee salary is higher than at least one salary in dept 10.
WHERE salary > ANY (
    SELECT salary FROM Employee WHERE dno = 10
);
```

Summary

- **Basic Subquery** - >Query inside query (ex. Compare salary with Ahmed Ali)
- **ALL** - > Must satisfy condition vs **all values** (ex. Salary > all salaries in dept 10)
- **ANY / SOME** - > Must satisfy condition vs **at least one value** (ex. Salary > any salary in dept 10)

SQL Aggregate Functions

What Are Aggregate Functions?

- Perform a calculation on **a set of values** → return **1 value**.
- Ignore **NULLs** (except **COUNT(*)**).

Common aggregate functions:

- **MAX()** → highest value
- **MIN()** → lowest value
- **AVG()** → average value
- **SUM()** → total sum
- **COUNT()** → count of rows (ignores NULLs)

Tip: Always use aliases (**AS**) for clarity.

- - Highest & lowest salary
`SELECT MAX(salary) AS max_salary,
MIN(salary) AS min_salary FROM Employee;`
- - Average salary
`SELECT AVG(salary) AS avg_salary FROM Employee;`
- - Total salary
`SELECT SUM(salary) AS total_salary FROM Employee;`
- - Count employees with non-null salary
`SELECT COUNT(salary) AS employee_count FROM Employee;`

- - Count all rows (even if salary is NULL)
SELECT COUNT(*) AS total_employees FROM Employee;

GROUP BY & HAVING

- **GROUP BY** → groups rows by column(s).
- **HAVING** → filters **groups** after aggregation.
- **WHERE** → filters **rows** before grouping.

Example: Average Salary per Department

```
SELECT department_number, AVG(salary) AS avg_salary FROM Employee
GROUP BY department_number HAVING MAX(salary) > 1800;
-- Show only departments where the top salary > 1800
```

SQL Views

What is a View?

- A **View** = **logical table** (virtual table).
- **Does not store data** → only stores the query definition.
- Acts like a **window** to underlying tables.
- Uses:
 - Simplify complex queries
 - Restrict sensitive data
 - Present data in a customized way

Creating a View

```
CREATE VIEW EmployeeProject AS
SELECT e.fname,e.lname,p.pname,w.hours FROM Employee e
JOIN WorksOn w ON e.emp_id = w.emp_id
JOIN Project p ON w.project_id = p.project_id;
```

- View Name → EmployeeProject
- Combines data from multiple tables (Employee, WorksOn, Project)

Querying a View

```
SELECT * FROM EmployeeProject -- Query a View just like a regular table.  
WHERE hours > 20;
```

Updating / Modifying Views

- **Simple View** (single table, no joins/aggregates) → DML (**INSERT** , **UPDATE** , **DELETE**) usually **allowed**.
- **Complex View** (joins, aggregates, subqueries) → DML often **restricted**.
- Modify View:

```
CREATE OR REPLACE VIEW EmployeeProject AS -- new definition here
```

Delete a view

```
DROP VIEW EmployeeProject;
```

Advanced Options

- **WITH CHECK OPTION** → Ensures any **INSERT/UPDATE** via the View satisfies the View condition.

```
CREATE VIEW HighSalaryEmployees AS  
SELECT fname, lname, salary FROM Employee  
WHERE salary > 3000 WITH CHECK OPTION;
```

- **Access control** → Grant permissions on the View instead of base table:

```
GRANT SELECT ON HighSalaryEmployees TO UserX;
```

Advantages of Views

1. **Simplifies queries** (hide joins/logic).
2. **Restricts access** (security & privacy).
3. **Consistency** (reuse same query definition).
4. **Abstraction** (users don't need to know underlying schema).

SQL Index

What is an Index?

- **Index** = Database object that **speeds up data retrieval**.
- Works like a **phonebook** → find rows quickly **without scanning the whole table**.
- Built on **one or more columns**, stores **sorted values + row pointers**.

Advantages

- Much faster **SELECT** queries.
- Improves performance of **JOIN**, **WHERE**, and **ORDER BY**.

Disadvantages

- Slows down **INSERT** / **UPDATE** / **DELETE** (because index must be updated).
- Consumes extra **storage space**.

Best Practices

- Index **columns used often in** **WHERE**, **JOIN**, **ORDER BY**, **GROUP BY**.
- Avoid indexing columns that:
 - Change frequently
 - Have **low selectivity** (e.g., gender **M/F**)
- Syntax & Examples:
 - 1. Create index on last name
`CREATE INDEX idx_employee_name ON Employee (lname);`
 - 2. Composite index (last name + department)
`CREATE INDEX idx_emp_lname_dept ON Employee (lname, department_id);`
 - 3. Drop index
`DROP INDEX idx_employee_name;`
 - Note: The DBMS decides automatically when to use an index in a query.

Key Idea

- **Indexes = Speed for SELECT, Cost for DML (Insert/Update/Delete).**
- Use wisely: balance **performance vs overhead.**